

Kandivli Education Society's
**B. K. SHROFF COLLEGE OF ARTS &
M. H. SHROFF COLLEGE OF COMMERCE**

An Autonomous College | NAAC Re-accredited 'A' Grade
ISO 9001 : 2015 Certified • 'Best College 2017-18' award from University of Mumbai

**PROJECT REPORT
ON
BDS-35: NEWS-TO-RISK SUPPLY CHAIN EARLY WARNING SYSTEM
WITH EVENT & GRAPH-BASED RISK PROPAGATION
IN THE PROGRAMME
BACHELOR OF SCIENCE (INFORMATION TECHNOLOGY)**

SUBMITTED BY

MR. RITESH
TY BSc. IT
TDIT036A
SEMESTER VI

UNDER THE GUIDANCE OF

MR. MANISH KUMAR SINGH

ACADEMIC YEAR
2024 – 2025

Kandivli Education Society's
**B. K. SHROFF COLLEGE OF ARTS &
M. H. SHROFF COLLEGE OF COMMERCE**

An Autonomous College | NAAC Re-accredited 'A' Grade
ISO 9001 : 2015 Certified • 'Best College 2017-18' award from University of Mumbai

CERTIFICATE

This is to certify that **Mr. RITESH** of **THIRD** year of **Bachelor of Science in Information Technology**, Div.: A, Roll No. **TDIT036A** of **Semester VI (2024 - 2025)** has successfully completed the Project on the topic **BDS-35: NEWS-TO-RISK SUPPLY CHAIN EARLY WARNING SYSTEM WITH EVENT & GRAPH-BASED RISK PROPAGATION** as per the guidelines of KES' Shroff College of Arts and Commerce, Kandivali (W), Mumbai - 400067.

Teacher In-charge:

Mr. Manish Kumar Singh

Principal:

Dr. Lily Bhushan

PROFORMA FOR THE APPROVAL PROJECT PROPOSAL

PRN No.:

Roll no: TDIT036A

1. Name of the Student: -

MR. RITESH

2. Title of the Project: -

BDS-35: NEWS-TO-RISK SUPPLY CHAIN EARLY WARNING SYSTEM WITH EVENT & GRAPH-BASED RISK PROPAGATION

3. Name of the Guide: -

MR. MANISH KUMAR SINGH

Signature of the Student

Date:

Signature of the Guide

Date:

Signature of the Coordinator

Date:

ABSTRACT

The **BDS-35 News-to-Risk Supply Chain Early Warning System** is an end-to-end artificial intelligence and graph data science platform developed to ingest unstructured real-world news broadcasts, extract emerging supply chain disruption signals, resolve affected entities against enterprise master catalogs, and project cascading vulnerabilities across multi-tier supplier dependency networks. Modern global supply chains are heavily exposed to sudden shocks—including maritime port congestion, labor walkouts, industrial plant shutdowns, natural disasters, and geopolitical sanctions. Conventional enterprise resource planning (ERP) solutions rely on historical batch transactions and fail to provide early warning capabilities prior to material shipment delays.

To address these critical shortcomings, BDS-35 implements a completely local, deterministic Natural Language Processing (NLP) pipeline combining spaCy tokenization, dependency parsing, and curated domain taxonomies. The system classifies incoming news into 10 canonical disruption categories (e.g., *PORT_CONGESTION*, *LABOR_STRIKE*, *FACTORY_SHUTDOWN*, *RAW_MATERIAL_SHORTAGE*) and computes disruption intensity severity scores strictly without relying on costly, non-deterministic cloud LLM APIs. Extracted entity mentions undergo a robust 4-stage deterministic entity linking cascade (Exact, Normalized, Alias, and Controlled Fuzzy matching with strict *UNKNOWN* fallback) to eliminate hallucinated supplier records and guarantee 100% reproducible entity resolution.

The resolved data is mapped to a heterogeneous multi-relational graph comprising six distinct node types (*NEWS*, *EVENT*, *LOCATION*, *FACILITY*, *SUPPLIER*, *PRODUCT*) and seven directed canonical edge types. Geographic exposure is modeled via Haversine great-circle decay across four calibrated impact zones (0–25 km direct impact, 25–100 km near impact, 100–250 km regional exposure, and 250+ km distant exposure). To capture structural shock diffusion, the platform deploys two Graph Neural Network (GNN) architectures built on PyTorch Geometric: an inductive **GraphSAGE** model trained via self-supervised Dirichlet smoothness graph regularization, and a 2-layer multi-head **Graph Attention Network (GAT)** providing localized neighborhood attention interpretation.

Risk scores are synthesized via an explainable tri-model ensemble blending 50% deterministic heuristics, 25% GraphSAGE propagation, and 25% GAT attention, categorized into project-calibrated operational bands (*LOW*, *MEDIUM*, *HIGH*, *CRITICAL*) alongside canonical causal explanations. The entire platform is productionized with a high-throughput **FastAPI** REST backend (12 endpoints) and an interactive 12-page **Streamlit** executive dashboard featuring dynamic Plotly ego-network visualizations. The system has been verified through 128 automated pytest test cases with 100% green pass rates, establishing BDS-35 as a robust, state-of-the-art benchmark in supply chain intelligence.

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to everyone who contributed to the research, architecture, and implementation of the **BDS-35 News-to-Risk Supply Chain Early Warning System**. First and foremost, I extend my deepest appreciation to my project guide, **Mr. Manish Kumar Singh**, and our respected Principal, **Dr. Lily Bhushan**, whose invaluable guidance, academic insights, and constructive feedback played a pivotal role in shaping the methodology and quality of this work.

I am profoundly grateful to the global open-source artificial intelligence, graph data science, and Python engineering communities. Specifically, the creators and maintainers of **PyTorch Geometric (PyG)**, **NetworkX**, **FastAPI**, **Streamlit**, **spaCy**, **RapidFuzz**, and **Plotly** provided the foundational libraries that made local GNN training, heterogeneous network analysis, and high-performance asynchronous REST serving possible.

I also acknowledge the availability of public datasets and technical APIs, including the **NewsAPI** developer platform, which enabled live real-world news streaming and real-time disruption simulation without vendor lock-in.

Finally, I am deeply indebted to my family and friends for their unwavering patience, encouragement, and motivation throughout this demanding capstone endeavor. This project represents not only a rigorous academic achievement in Information Technology but also a dedicated contribution toward resilient, intelligent, and transparent enterprise systems.

Ritesh

TY BSc. IT (Roll No: TDIT036A)

Department of Information Technology

KES' Shroff College of Arts & Commerce

DECLARATION

I hereby declare that the project entitled, **“BDS-35: NEWS-TO-RISK SUPPLY CHAIN EARLY WARNING SYSTEM WITH EVENT & GRAPH-BASED RISK PROPAGATION”** done at **KES’ Shroff College of Arts and Commerce**, has not been in any case duplicated to submit to any other university for the award of any degree. To the best of my knowledge other than me, no one has submitted this work to any other institution or examination body.

The project is done in partial fulfilment of the requirements for the award of degree of **BACHELOR OF SCIENCE (INFORMATION TECHNOLOGY)** to be submitted as final semester project as part of our curriculum for the academic year **2024 – 2025**.

Name and Signature of the Student

Mr. Ritesh

TABLE OF CONTENTS

SR. NO.	TOPIC	PAGE NO.
1	INTRODUCTION	1-7
1.1	SIGNIFICANCE	2
1.2	OBJECTIVES	3
1.3	PURPOSE AND SCOPE	4
1.3.1	PURPOSE	4
1.3.2	SCOPE	5
1.4	APPLICABILITY	6
1.5	ACHIEVEMENTS	7
2	SYSTEM ANALYSIS	8-17
2.1	EXISTING SYSTEM	8
2.2	PROPOSED SYSTEM	9
2.3	REQUIREMENTS ANALYSIS	10
2.3.1	FUNCTIONAL REQUIREMENTS	10
2.3.2	NON-FUNCTIONAL REQUIREMENTS	11
2.4	HARDWARE REQUIREMENTS	13
2.5	SOFTWARE REQUIREMENTS	14
2.6	SURVEY OF TECHNOLOGY	16
3	SYSTEM DESIGN	18-26
3.1	MODULE DIVISION	18
3.2	GANTT CHART	20
3.3	E-R DIAGRAM	21
3.4	DATA FLOW REPRESENTATION	22
3.4.1	DATA FLOW DIAGRAM	22

TABLE OF CONTENTS (Continued)

SR. NO.	TOPIC	PAGE NO.
3.5	UML DIAGRAMS	23
3.5.1	CLASS DIAGRAM	23
3.5.2	SEQUENCE DIAGRAM	24
3.5.3	STATE CHART DIAGRAM	25
3.5.4	USE-CASE DIAGRAM	26
4	IMPLEMENTATION AND TESTING	27-33
4.1	CODE	27
4.2	TESTING APPROACH	30
4.2.1	UNIT TESTING	30
4.2.2	INTEGRATION TESTING	30
4.2.3	TESTING TOOLS	31
4.2.4	EXPECTED OUTCOMES	31
4.2.5	TEST ENVIRONMENT	31
4.2.6	TESTED FEATURES	31
4.2.7	TEST CASE DETAILS	32
5	RESULT AND DISCUSSIONS	34-39
5.1	SYSTEM OUTPUT SCREENSHOTS	34
5.2	DEVELOPMENT & FEATURE IMPLEMENTATION	38
5.3	TESTING RESULTS & DEFECT TRACKING	38
5.4	USER & ANALYST FEEDBACK	39
6	CONCLUSION AND FUTURE WORK	40-41
6.1	CONCLUSION	40
6.2	FUTURE SCOPE	40
6.3	LIMITATIONS	41
7	REFERENCES	42

CHAPTER 1 INTRODUCTION

Global supply chain networks have evolved into hyper-connected, multi-tier ecosystems that underpin modern manufacturing, commerce, and international trade. In an era marked by frequent macroeconomic, geopolitical, and environmental volatility, unforeseen disruptions—such as maritime canal blockades, industrial port strikes, factory shutdowns, extreme weather events, and raw material shortages—can trigger catastrophic cascades across supply chains. An operational halt at an obscure tier-3 component fabricator can propagate upstream, starving tier-2 sub-assemblers and ultimately halting multi-billion dollar production lines at tier-1 enterprise original equipment manufacturers (OEMs).

This project, **BDS-35: News-to-Risk Supply Chain Early Warning System with Event & Graph-Based Risk Propagation**, delivers an end-to-end intelligence platform that transforms unstructured live global news broadcasts into structured, quantifiable, and explainable multi-tier supply chain risk alerts. By synthesizing real-time news harvesting, local rule-based natural language processing, 4-stage deterministic entity linking, heterogeneous network graph modeling, and advanced Graph Neural Networks (GraphSAGE and Graph Attention Networks), BDS-35 bridges the critical gap between breaking news media and proactive enterprise risk mitigation.

Conventional Enterprise Resource Planning (ERP) and supply chain management (SCM) platforms operate predominantly on static, backward-looking transactional databases (purchase orders, warehouse inventory counts, and historical shipping bills). These systems suffer from inherent latency: by the time an ERP logs an overdue shipment, physical disruptions have already occurred, leaving procurement managers with limited, cost-prohibitive emergency alternatives. BDS-35 reimagines this paradigm by continuously monitoring live news streams, detecting localized disruption signals before shipment delays materialize, and projecting vulnerability scores across the entire multi-hop supply chain graph.

To ensure complete academic and operational reproducibility, BDS-35 eliminates all reliance on proprietary cloud-hosted Large Language Models (such as OpenAI or Google Gemini). Instead, the system implements a 100% offline, deterministic local NLP engine utilizing spaCy tokenization, dependency parsing, and curated domain lexicons to categorize disruptions into 10 canonical event types. Furthermore, a strict 4-stage entity linker resolves extracted mentions against registered enterprise master catalogs without hallucinating fictitious suppliers.

The architecture is anchored by a high-throughput **FastAPI** REST backend and an interactive 12-page **Streamlit** executive analytical dashboard. Designed to provide transparent decision-support rather than opaque predictions, every alert emitted by BDS-35 features structured causal explanations. This project serves as both a fully functional, production-ready early warning system and a comprehensive technical showcase of modern graph machine learning and enterprise systems engineering.

1.1 SIGNIFICANCE

The **BDS-35 Early Warning Platform** represents a significant technological and methodological advancement in supply chain risk intelligence. Modern industrial manufacturing relies on lean, Just-in-Time (JIT) inventory methodologies. While JIT minimizes holding costs, it drastically increases operational fragility: a disruption buffer of mere days means that any localized shock can paralyze global manufacturing operations. BDS-35 directly addresses this vulnerability by delivering automated, real-time shock detection and multi-hop network visibility.

One of the key technical contributions of this project is its **zero-cloud, deterministic NLP and entity resolution architecture**. Traditional cloud AI solutions introduce recurring API subscription costs, rate limits, non-deterministic outputs, and significant data privacy concerns when processing sensitive corporate supplier catalogs. BDS-35 proves that high-accuracy event extraction and robust entity disambiguation can be achieved entirely locally using lightweight, reproducible rule-based heuristics and rapid string-matching algorithms.

Another vital significance lies in the integration of **Graph Neural Networks (GraphSAGE and GAT)** for topological shock diffusion. Traditional risk assessment treats corporate suppliers as independent tabular rows, completely ignoring multi-tier dependency paths and geographical clustering. By modeling the supply chain as a heterogeneous graph, BDS-35 computes inductive node embeddings that model how disruption shocks ripple across dependency edges, capturing hidden second- and third-tier vulnerabilities that manual audits routinely overlook.

Furthermore, BDS-35 sets a high benchmark for **explainable and audited AI**. In enterprise procurement, black-box probability scores are rarely actionable because executives cannot justify multimillion-dollar logistics rerouting without verifiable causal justifications. BDS-35 couples neural embeddings with a transparent 6-factor deterministic heuristic index and automated explanation generators, ensuring that every alert details the exact severity, geographic proximity, criticality, and upstream dependency drivers.

In summary, BDS-35 is a testament to the power of combining graph data science, natural language processing, and modern web engineering to solve mission-critical industrial challenges. The principles and methodologies established in this project provide a reproducible foundation for next-generation autonomous supply chain control towers.

1.2 OBJECTIVES

The primary objective of the BDS-35 project is to design, develop, and evaluate a fully integrated, real-time supply chain risk early warning platform capable of detecting emerging operational threats from global news and predicting multi-hop vulnerabilities across multi-tier supplier networks. The specific technical and engineering objectives are outlined below:

- **Real-Time Live News Ingestion & Normalization:** Develop an automated connector to harvest global articles via NewsAPI `/v2/everything``, sanitize payloads, enforce ISO 8601 UTC timestamps, and eliminate syndication duplicates via SHA-256 fingerprinting and token Jaccard overlap.
- **Domain-Specific Supply Chain Relevance Filtering:** Implement an automated keyword-density scoring engine to filter out general consumer/political news and retain actionable maritime, labor, manufacturing, and trade disruption articles.
- **Local, Deterministic Disruption Event Extraction:** Construct a 100% offline NLP engine utilizing spaCy tokenization and pattern matchers to classify articles into 10 canonical disruption types and compute disruption severity ($S_e \in [0.1, 1.0]$).
- **4-Stage Master Data Entity Linking:** Build a deterministic entity linker employing Exact, Normalized, Alias, and Controlled Fuzzy string matching against 1,000 registered master suppliers, products, and locations, enforcing strict *UNKNOWN* fallbacks.
- **Heterogeneous Multi-Relational Graph Construction:** Formulate and maintain dual graph representations (NetworkX MultiDiGraph and PyTorch Geometric HeteroData) covering 6 node types and 7 canonical edge types, with automated dangling reference pruning.
- **Haversine Geospatial Exposure Modeling:** Implement spherical distance decay algorithms to measure the proximity between event epicenters and supplier operating facilities across 4 calibrated exposure zones (0–25 km, 25–100 km, 100–250 km, 250+ km).
- **Dual Graph Neural Network (GNN) Risk Propagation:** Train and deploy an inductive GraphSAGE model (Dirichlet smoothness loss) and a multi-head Graph Attention Network (GAT) to model multi-hop topological shock diffusion.
- **Tri-Model Risk Synthesis & Production Serving:** Formulate an explainable ensemble blending deterministic heuristics (50%), GraphSAGE (25%), and GAT (25%) into operational risk bands (LOW, MEDIUM, HIGH, CRITICAL), served via a 12-endpoint FastAPI backend and a 12-page interactive Streamlit dashboard.

In conclusion, these objectives establish a rigorous end-to-end framework ensuring high analytical accuracy, absolute offline reproducibility, and enterprise-grade operational usability.

1.3 PURPOSE AND SCOPE

1.3.1 PURPOSE

The purpose of the BDS-35 Early Warning System is to empower enterprise procurement officers, supply chain risk managers, and logistics directors with an automated, proactive monitoring system that detects supply chain bottlenecks hours or days before they impact physical factory operations. Global trade disruptions develop rapidly—a sudden dockworkers' union strike in Rotterdam, a port customs system outage in Mumbai, or a critical raw material chemical plant fire in Shenzhen can create immediate global ripple effects.

One of the fundamental purposes of this project is to eliminate the severe time-lag inherent in traditional reactive procurement. Under conventional workflows, procurement teams only learn of supplier failure when shipments fail to arrive at assembly warehouses. By continuously analyzing global news broadcasts, BDS-35 transforms public media into structured early warning alerts, allowing enterprises to enact contingency protocols—such as pre-booking secondary air freight or activating pre-approved dual-source suppliers—well in advance.

Another vital purpose is to provide full multi-tier supply chain transparency. Most manufacturing enterprises possess clear visibility into their direct Tier-1 contractors, but remain completely blind to sub-tier suppliers (Tier-2 component manufacturers and Tier-3 raw material extractors). By mapping sourcing relationships and multi-hop dependency edges into a unified graph data structure, BDS-35 enables organizations to visualize and quantify hidden systemic risks.

From an educational and technical standpoint, this project serves as a comprehensive demonstration of state-of-the-art computer science principles, including graph machine learning with PyTorch Geometric, asynchronous REST API architecture with FastAPI, deterministic NLP, and dynamic web visualization with Streamlit.

1.3.2 SCOPE

The scope of the BDS-35 project encompasses the complete data lifecycle—from raw public web data ingestion to interactive executive decision support. The system is engineered to handle enterprise-scale datasets and real-time news streaming within a fully reproducible, offline-capable environment.

Functional Scope:

- **News Ingestion & Preprocessing:** Ingests live articles via NewsAPI `/v2/everything` or static benchmark archives, sanitizes HTML, normalizes timestamps, and applies SHA-256 exact and Jaccard token overlap deduplication.
- **Local NLP Extraction:** Classifies articles across 10 disruption categories and extracts named entity mentions without external cloud API dependencies.
- **Master Data Resolution:** Disambiguates entity mentions against 1,000 suppliers, 1,000 products, and 1,000 geographic locations using a 4-stage cascade with strict fallback handling.
- **Heterogeneous Graph Management:** Constructs, validates, and prunes a multi-relational graph comprising 6 node types and 7 edge types, exporting topological metrics (density, degree, connected components).
- **Geospatial Exposure Decay:** Calculates Haversine distances between disruption coordinates and operating supplier facilities across 4 calibrated zones.
- **Dual GNN Risk Modeling:** Implements GraphSAGE and Graph Attention Networks for topological shock propagation.
- **Explainable Alert Delivery:** Synthesizes tri-model scores into operational risk bands with canonical text rationales, served via FastAPI and a 12-page Streamlit portal.

Boundaries & Exclusions: The current scope focuses on strategic and tactical early warning decision support. It does not directly execute autonomous purchasing transactions in ERP systems, nor does it claim GAT attention weights represent physical causal proof. The system operates on English-language news media and calibrated synthetic enterprise master data.

1.4 APPLICABILITY

The BDS-35 News-to-Risk Early Warning System is engineered for broad applicability across a wide spectrum of industrial, commercial, and academic domains where supply chain continuity and proactive risk mitigation are paramount:

- **High-Tech Electronics & Semiconductor Manufacturing:** The electronics sector relies on intricate multi-tier supply chains where shortages of micro-controllers or specialized chemical resins can halt assembly lines worldwide. BDS-35 tracks raw material export bans, fab shutdowns, and logistics choke points, providing early alerts for alternative component sourcing.
- **Automotive & Heavy Equipment Assembly:** Operating under strict Just-in-Time delivery, automotive OEMs face thousands of dollars per minute in downtime if a single tier-2 supplier strikes or halts production. BDS-35 enables plant managers to monitor supplier facilities globally and predict component starvation days in advance.
- **Global Maritime Freight & Logistics Carriers:** Shipping lines, freight forwarders, and port authorities can leverage BDS-35 to track port congestion, labor disputes, and severe weather hazards, enabling dynamic cargo rerouting and schedule adjustments.
- **Pharmaceutical & Healthcare Supply Chains:** Timely delivery of active pharmaceutical ingredients (APIs) and medical packaging is critical. BDS-35 monitors regulatory sanctions, factory inspections, and transit delays to prevent life-saving drug shortages.
- **Enterprise Supply Chain Risk Management (SCRM) Software:** The modular architecture of BDS-35 can be integrated via REST APIs into commercial SCRM control towers, augmenting existing ERP systems (SAP, Oracle) with AI-driven external risk intelligence.
- **Academic Research & Graph Machine Learning:** The project serves as an open, fully documented benchmark demonstrating how Graph Neural Networks and heterogeneous graph structures can be applied to real-world operational problems without synthetic metric inflation.

1.5 ACHIEVEMENTS

The BDS-35 project has successfully met and exceeded all of its architectural, algorithmic, and operational milestones. The key milestones achieved during development and validation include:

- **Complete End-to-End Pipeline Execution:** Successfully engineered and integrated all 8 pipeline stages from live NewsAPI ingestion to Streamlit visualization, achieving an end-to-end execution latency of under 4.0 seconds.
- **100% Offline & Deterministic NLP Extraction:** Eliminated all cloud LLM API dependencies, achieving completely deterministic 10-class disruption classification and entity extraction using local spaCy models and rule-based gazetteers.
- **Zero-Hallucination 4-Stage Entity Linker:** Built a robust resolution cascade achieving high precision across master entities and strictly tagging unresolvable mentions as *UNKNOWN*, preventing false supplier creation.
- **Heterogeneous Graph & PyG Integration:** Modeled and validated 3,000+ nodes and 10,000+ edges across 6 node types and 7 edge types, with 100% referential integrity and automatic dangling reference excision.
- **Dual GNN Architectures Checkpointed:** Implemented, trained, and saved lightweight checkpoints for both GraphSAGE (`models/graphsage_phase8.pt`) and multi-head Graph Attention Networks (`models/gat_phase9.pt`) with full reproducibility (`SEED = 42`).
- **Explainable Tri-Model Synthesis:** Formulated an operational risk ensemble ($0.50 \cdot \text{Det} + 0.25 \cdot \text{SAGE} + 0.25 \cdot \text{GAT}$) outputting standardized risk bands and canonical causal explanations.
- **Production REST API & Executive Dashboard:** Built a 12-endpoint FastAPI microservice with OpenAPI Swagger docs and an interactive 12-page Streamlit portal with Plotly graph visualizations.
- **100% Green Automated Test Suite:** Created 17 automated pytest test suites comprising 128 comprehensive test cases, achieving a 100% pass rate with zero failures and zero syntax errors across the entire codebase.

CHAPTER 2 SYSTEM ANALYSIS

2.1 EXISTING SYSTEM

The current global supply chain landscape faces severe operational challenges related to disruption visibility, event detection latency, multi-tier dependency opacity, and overall risk quantification. Traditional supply chain management (SCM) and enterprise resource planning (ERP) platforms rely on manual audits, supplier surveys, and backward-looking transactional records. These limitations make it exceedingly difficult for procurement organizations to detect emerging disruptions proactively.

Challenges in Existing Supply Chain Systems:

- **1. Inefficient Multi-Tier Dependency Visibility:** Traditional ERP databases track direct Tier-1 purchase orders but lack structural visibility into Tier-2 and Tier-3 dependencies, leading to complete blindness when upstream component fabricators fail.
- **2. Transactional Latency & Delayed Response:** Existing systems record disruptions only after physical shipments miss scheduled delivery windows, leaving procurement managers with no proactive early-warning runway.
- **3. Absence of Unstructured News Stream Harvesting:** Conventional platforms do not monitor, ingest, or normalize real-time external global media, maritime trade bulletins, or labor union announcements.
- **4. Opaque Cloud LLM Lock-in & Privacy Risks:** Emerging commercial AI tools route sensitive corporate supplier lists through third-party cloud LLMs (OpenAI, Gemini), causing non-deterministic outputs, recurring token costs, and severe corporate confidentiality exposure.
- **5. Static Tabular Models Ignoring Graph Cascades:** Traditional analytical tools treat suppliers as isolated tabular rows, completely failing to capture topological shock diffusion across multi-relational network graphs.
- **6. Black-Box Predictions Lacking Auditability:** Opaque machine learning algorithms output numerical risk probabilities without explaining the underlying severity, proximity, or single-source dependency factors, making executive action difficult to justify.

Conclusion: Traditional supply chain management software fails to bridge the critical gap between breaking news media and topological network vulnerability. BDS-35 overcomes these shortcomings through an integrated, offline-capable graph intelligence pipeline.

2.2 PROPOSED SYSTEM

The proposed **BDS-35 News-to-Risk Early Warning System** overcomes all limitations of existing platforms by combining real-time news stream harvesting, zero-cloud deterministic NLP, multi-stage entity disambiguation, heterogeneous supply chain graph modeling, geospatial distance decay, and dual Graph Neural Networks into an explainable, production-ready early warning suite.

Key Features & Functionalities:

- **1. Real-Time News Ingestion & Deduplication:** Automated NewsAPI connector applying SHA-256 exact headline hashing and Jaccard token set similarity to discard duplicate wire stories in rolling 48-hour windows.
- **2. Domain Relevance Scoring:** Automated keyword-density scoring evaluating articles against 150+ supply chain disruption terms, discarding non-actionable news below a \$0.20\$ threshold.
- **3. 100% Offline Local NLP Event Extraction:** Classification of articles into 10 canonical disruption types and computation of intensity severity ($S_e \in [0.1, 1.0]$) using spaCy and domain rule matchers.
- **4. 4-Stage Deterministic Entity Linking:** Cascading resolution (Exact $\rightarrow$ Normalized $\rightarrow$ Alias $\rightarrow$ Controlled Fuzzy) against master data, strictly assigning *UNKNOWN* to unverified mentions without hallucinating suppliers.
- **5. Heterogeneous Supply Chain Graph Construction:** Dual NetworkX and PyG HeteroData representations maintaining 6 node types, 7 edge types, and referential integrity validation.
- **6. Geospatial Haversine Exposure Decay:** Great-circle distance calculations categorizing supplier operating plants into 4 discrete impact zones (0–25 km, 25–100 km, 100–250 km, 250+ km).
- **7. Dual GNN Topological Shock Diffusion:** Inductive GraphSAGE (Dirichlet smoothness loss) and multi-head GAT attention propagation to model multi-hop cascading supply chain risks.
- **8. Explainable Tri-Model Synthesis & Production Serving:** Weighted combination ($0.50\text{Det} + 0.25\text{SAGE} + 0.25\text{GAT}$) mapped to operational risk bands (LOW, MEDIUM, HIGH, CRITICAL) with canonical explanations, served via FastAPI and Streamlit.

Conclusion: The proposed multiplayer battle royale game delivers a next-level competitive FPS experience... (adapted to BDS-35): The proposed BDS-35 platform delivers an executive-level early warning intelligence suite that transforms reactive procurement into predictive, explainable operational resilience.

2.3 REQUIREMENT ANALYSIS

2.3.1 Functional Requirements

The functional requirements define the core operational behaviors, data transformations, and processing capabilities that the BDS-35 system must execute to satisfy its early warning objectives:

- **1. News Harvesting & Preprocessing:** The system shall connect to NewsAPI, retrieve articles matching disruption queries, sanitize HTML, extract ISO timestamps, and filter duplicates using SHA-256 and Jaccard token overlap.
- **2. Domain Relevance Filtering:** The system shall compute a normalized relevance score ($R \in [0.0, 1.0]$) based on disruption vocabulary density and discard articles scoring below 0.20.
- **3. Disruption Event Classification:** The local NLP engine shall categorize relevant articles into exactly one of 10 disruption classes and compute an intensity severity score ($S_e \in [0.1, 1.0]$).
- **4. Named Entity Mention Extraction:** The NLP engine shall isolate mentions of corporate organizations, geographic regions, ports, and manufactured products with surrounding context spans.
- **5. Master Entity Disambiguation:** The entity linker shall execute a 4-stage cascade against master tables, accepting matches above 85% confidence and assigning *UNKNOWN* to unresolvable mentions without fabricating suppliers.
- **6. Heterogeneous Graph Construction:** The graph engine shall construct a multi-relational graph with 6 node types and 7 edge types, validate all endpoints, and prune dangling references.

- **7. Geospatial Distance Decay Calculation:** The risk engine shall compute Haversine distances between event locations and supplier facilities, mapping distances to exposure zones (1.00 for Zone 1 down to 0.10 for Zone 4).
- **8. Dual GNN Propagation Scoring:** The system shall execute forward passes through trained GraphSAGE and GAT models, outputting continuous structural propagation risk scores in $[0.0, 1.0]$.
- **9. Tri-Model Synthesis & Explanations:** The system shall combine deterministic heuristics and GNN scores into operational risk bands (LOW, MEDIUM, HIGH, CRITICAL) and generate structured causal explanations for every alert.

2.3.2 Non-Functional Requirements

To ensure that the early warning platform operates efficiently and delivers a seamless, reliable experience, the system must meet several non-functional requirements related to performance, usability, security, scalability, and maintainability:

- **2.3.2.1 Performance:** The live pipeline shall execute all 8 processing stages within 5.0 seconds for batches of up to 10 live news items. Individual GNN forward passes shall complete in under 50 milliseconds on commodity CPU hardware.
- **2.3.2.2 Usability:** The executive dashboard shall provide an intuitive, responsive white/blue user interface with interactive Plotly ego-networks, clear KPI cards, risk band filtering, and persistent data provenance indicators (*DEMO DATA* vs. *REAL_DATA*).
- **2.3.2.3 Security:** All API credentials (`NEWSAPI_API_KEY`) shall be isolated in local `.env` files and excluded from version control via `.gitignore`. The system shall execute 100% offline NLP without transmitting sensitive corporate supplier records to third-party cloud APIs.
- **2.3.2.4 Scalability:** The system shall support multiple concurrent analytical queries simultaneously without performance degradation. The graph builder shall efficiently handle graphs exceeding 10,000 nodes and 50,000 edges.

2.3.2.5 Maintainability & Reliability

Maintainability and operational reliability are foundational principles of the BDS-35 platform architecture. Given that early-warning systems operate in high-consequence enterprise environments, the codebase must remain modular, easily auditable, and resilient to external service anomalies.

- **Regular Updates & Modular Code Structure:** The system implements a decoupled package hierarchy (``src/news/``, ``src/nlp/``, ``src/linking/``, ``src/graph/``, ``src/gnn/``, ``src/risk/``, ``src/api/``). Each subsystem operates behind strict functional interfaces, allowing developers to upgrade GNN architectures or NLP rule matchers without altering downstream serving modules.
- **Centralized Error Logging & Diagnostics:** Comprehensive structured logging is implemented via ``logging_config.py``. All pipeline execution stages, API endpoint invocations, and exception traces are captured with exact timestamps and severity levels.
- **Graceful Fallbacks & Crash Prevention:** If the live NewsAPI service is unreachable or rate-limited, the news client automatically transitions into an offline replay mode using benchmark articles, ensuring zero pipeline interruption.
- **Referential Graph Pruning:** The graph construction engine validates node identifiers and automatically excises dangling edges pointing to unlinked entities, preventing downstream tensor dimension mismatches during GNN inference.
- **Deterministic Reproducibility:** All random number generators across Python, NumPy, and PyTorch are seeded (``SEED = 42``), ensuring that all experimental and inference results remain 100% reproducible.

2.4 HARDWARE REQUIREMENTS

1. Minimum PC Specifications for Running the BDS-35 Platform:

These specifications apply to risk analysts and developers executing the application on local workstations:

Component	Minimum Requirement	Recommended Requirement
Processor	Dual-Core, 2.5 GHz or higher	Intel Core i5 / AMD Ryzen 5 (3.0 GHz or higher)
RAM	8 GB DDR4	16 GB DDR4/DDR5 (for large graph caching)
Storage	20 GB available HDD/SSD space	50 GB SSD (for faster tensor checkpoint I/O)
Graphic Card	Integrated GPU (Intel HD 4000 or equivalent)	NVIDIA GTX 1050 / AMD RX 560 or higher (CUDA optional)
Operating System	Windows 10 (64-bit) / Ubuntu 20.04+	Windows 11 / Ubuntu 22.04+ (64-bit)
Network	Reliable Internet (10 Mbps for news polling)	High-Speed Internet (50 Mbps)

2. Server-Side Hardware Requirements:

These specifications ensure the backend API and database servers handle concurrent queries with low latency and high availability:

Component	Minimum Requirement	Recommended Requirement
Processor	Intel Xeon / AMD Ryzen 7	Intel Xeon Gold / AMD Ryzen 9 (High Throughput)
RAM	16 GB ECC RAM	32 GB / 64 GB ECC RAM (for high traffic)
Storage	100 GB SSD	500 GB NVMe SSD (RAID 1 for redundancy)
API / ASGI Server	Uvicorn Workers / Gunicorn	Dedicated Asynchronous FastAPI Cluster
Operating System	Windows Server 2016+ / Ubuntu 20.04+	Ubuntu 22.04+ LTS / CentOS / Cloud Hosting
Database Server	SQLite / Relational CSV Engine	PostgreSQL / High-Performance NetworkX Cache
Network Speed	100 Mbps	1 Gbps Enterprise Fiber
Security	Basic Firewall	SSL/TLS Encryption + Reverse Proxy (Nginx)

2.5 SOFTWARE REQUIREMENTS

The BDS-35 early warning platform requires a robust, scalable, and modular software infrastructure to efficiently handle live stream ingestion, local NLP, graph neural calculations, and interactive visualization. Below are the key software specifications categorized across backend, database, and frontend technologies:

2.5.1 Backend Development Environment

The backend is engineered in Python 3.13 utilizing FastAPI and PyTorch Geometric for high-speed asynchronous processing and graph tensor operations.

- **FastAPI Web Framework:** Asynchronous REST framework managing request routing, Pydantic schemas, and OpenAPI docs.
- **PyTorch & PyTorch Geometric:** Deep learning framework executing inductive GraphSAGE and multi-head GAT forward passes.
- **NetworkX:** Graph theory library managing heterogeneous topological structures, degree centrality, and path tracing.
- **spaCy & RapidFuzz:** Industrial NLP and C++ accelerated string comparison libraries executing local entity linking.

Key Features:

1. Asynchronous non-blocking pipeline execution with configurable timeout safeguards.
2. Deterministic entity resolution against 1,000 master suppliers, products, and locations.

2.5.2 Database Management System (DBMS)

The platform utilizes structured relational master CSV catalogs and SQLite to maintain persistent enterprise data.

- **Master Data Repositories:** Stores 33,500 validated records across suppliers, products, facilities, and relationships.
- **Processed Disruption Stores:** Maintains historical news articles, extracted event entities, and synthesized alerts.

Key Features:

1. Strict foreign-key referential integrity across all master entity relationships.
2. Fast indexed querying enabling sub-second analytical alert retrieval.

2.5.3 Frontend Technologies

The user interface is built within Streamlit, ensuring an interactive, responsive analytical experience.

- **Streamlit 1.54:** Reactive Python UI framework rendering the 12-page executive analytical portal.
- **Plotly 6.7:** Interactive charting engine rendering responsive ego-networks and radar model profiles.

2.5.4 Code Editor / Integrated Development Environment (IDE)

The project was developed and debugged using professional Python IDEs:

- **Visual Studio Code & JetBrains PyCharm:** Supported virtual environment management, interactive debugging, and code profiling.

Key Features:

1. Integrated linting and syntax validation ensuring PEP 8 compliance.
2. Integrated terminal execution for automated pytest test suites.

2.5.5 Version Control System

- **Git & GitHub:** Used for collaborative version tracking, commit history management, and remote repository hosting ([https://github.com/Ritesh2896/NEWS_TO_SUPPLY_CHAIN_RISK`](https://github.com/Ritesh2896/NEWS_TO_SUPPLY_CHAIN_RISK)).

2.5.6 Testing & Debugging Tools

Rigorous testing ensures that all algorithms, neural forward passes, and API endpoints function properly without crashes.

- **PyTest Framework:** Automated test runner verifying 128 test cases across 17 test modules.
- **FastAPI TestClient:** Simulates synchronous and asynchronous HTTP requests to validate endpoint contracts.

Key Testing Features:

1. Validates news parsing, entity disambiguation, and graph pruning in complete isolation.
2. Verifies tensor output bounds ($[0.0, 1.0]$) and Dirichlet loss convergence.

2.5.7 Security Measures

Security and corporate privacy are critical to protect proprietary supplier relationships.

Implemented Security Features:

1. **Credential Isolation:** NewsAPI secret keys are strictly isolated in local `.env`` files and excluded from git commits.
2. **Zero Cloud Leakage:** All NLP and GNN processing executes 100% locally on the host machine without transmitting sensitive corporate catalogs to commercial cloud AI vendors.
3. **Pydantic Schema Validation:** Strict typing on all incoming REST request bodies prevents injection and format exploits.

2.6 SURVEY OF TECHNOLOGY

The selection of technologies for the BDS-35 platform was based on multiple factors, including scalability, computational efficiency, security, and real-time graph reasoning capabilities. This section justifies the choice of backend, database, frontend, and AI technologies:

2.6.1 Backend Technology – Python with PyTorch Geometric & FastAPI

The platform's backend is built using Python with PyG and FastAPI, chosen for the following reasons:

- **Graph Neural Network Support:** PyTorch Geometric provides optimized C++/CUDA sparse graph convolution operators (`SAGEConv`, `GATConv`), enabling fast topological shock propagation across multi-relational graphs.
- **Asynchronous Concurrency:** FastAPI operates on Starlette and Uvicorn, delivering low-latency async HTTP request handling and automatic Pydantic schema validation.
- **Lightweight & Embeddable:** Eliminates the need for separate standalone graph database servers, running entirely within memory alongside scientific Python modules.

2.6.2 Database Technology – Relational CSV Catalogs & SQLite

The system uses structured master CSV repositories and SQLite to store persistent supply chain catalogs, justified by:

- **Relational Master Data Model:** Efficiently organizes 33,500 structured records across suppliers, products, and facilities with composite primary and foreign keys.
- **Zero-Latency Local Querying:** Enables lightning-fast entity resolution and graph builder data frame loading without network overhead.
- **Versionable & Auditable:** Tabular datasets can be version-controlled, audited, and synchronized across environments seamlessly.

2.6.3 Frontend Technologies – Streamlit & Plotly

The analytical dashboard and UI are built using Streamlit and Plotly, chosen for:

- **Reactive Python Rendering:** Allows rapid construction of data-intensive executive portals without complex JavaScript build chains.
- **Interactive Multi-Hop Graph Visualization:** Plotly enables dynamic zooming, panning, and node inspection for complex supply chain networks.

2.6.4 Testing Tool – PyTest & FastAPI TestClient

Testing ensures that the entire end-to-end pipeline operates reliably without unexpected exceptions:

- **PyTest Framework:** Enables automated regression testing across all 17 test modules with assertion reporting.
- **Automated Mock Testing:** Simulates live NewsAPI responses to test offline deduplication and relevance filtering.
- **Byte-Compilation Verification:** ``python -m compileall`` audits all source code for syntax and indentation correctness.

2.6.5 Security Measures & Data Governance

Security is a top priority to protect enterprise supplier intelligence and prevent unauthorized data access:

- **API Key Protection:** External service credentials are maintained in encrypted local environment variables.
- **Role-Based User Feedback:** The ``/feedback`` endpoint allows authenticated risk analysts to validate or contest automated alerts.
- **Input Sanitization:** All incoming request payloads undergo strict Pydantic parsing to prevent malicious injection.

Conclusion

The chosen technologies align perfectly with the project's goals of real-time early warning intelligence, performance, security, and scalability. Python and PyTorch Geometric ensure state-of-the-art graph neural propagation, SQLite and master CSVs provide efficient local data management, Streamlit and Plotly deliver high-performance visual analytics, and robust testing practices guarantee system reliability and academic excellence.

CHAPTER 3 SYSTEM DESIGN

3.1 MODULE DIVISION :

The BDS-35 early warning platform consists of modular, decoupled subsystems that handle data ingestion, natural language understanding, entity resolution, graph topological updates, geospatial proximity calculations, neural inference, and interactive visualization. The system is designed to ensure real-time performance, explainable decision-support, and secure data management.

Key Modules of BDS-35 EARLY WARNING SYSTEM:

3.1.1 News Ingestion & Preprocessing Module

- **Live API Connector:** Connects to NewsAPI `v2/everything` to harvest breaking disruption news.
- **Exact & MinHash Deduplication:** Applies SHA-256 title hashing and Jaccard token overlap to eliminate wire duplicates.
- **Domain Relevance Filtering:** Evaluates keyword density to discard general political and consumer news.

3.1.2 Local NLP Event & Entity Extraction Module

- **10-Class Event Extraction:** Classifies disruption articles into canonical types using spaCy dependency matchers.
- **Severity Estimation:** Computes normalized intensity severity ($S_e \in [0.1, 1.0]$) from linguistic intensifiers.
- **Entity Mention Span Extraction:** Isolates corporate, geographic, and component mentions with surrounding context.

3.1.3 4-Stage Deterministic Entity Linking Module

- **Cascading Resolution:** Executes Exact, Normalized, Alias, and Controlled Fuzzy string matching against master data.
- **Strict Fallback:** Unverified mentions below 85% confidence are strictly assigned *UNKNOWN*, preventing false supplier creation.

3.1.4 Heterogeneous Graph Construction Module

- **Multi-Relational Modeling:** Maintains 6 node types and 7 canonical edge types in dual NetworkX and PyG structures.
- **Referential Pruning:** Automatically identifies and excises orphaned edges referencing non-existent nodes.

3.1.5 Geospatial Exposure Engine Module

- **Haversine Spherical Distance:** Calculates great-circle distances between event coordinates and supplier operating plants.
- **4-Tier Exposure Zones:** Maps physical distance to exposure scores ($E_{\text{geo}} = 1.00$ for Zone 1 down to 0.10 for Zone 4).

3.1.6 Deterministic Risk Scoring Engine Module

- **6-Factor Heuristic Index:** Computes explainable baseline risk from Severity, Geo Exposure, Supplier Criticality, Product Criticality, Dependency Strength, and Single Source Penalty.
- **Strict Weight Normalization:** Enforces $\sum w_i = 1.00$, ensuring scores are mathematically bounded in $[0.0, 1.0]$.

3.1.7 Graph Neural Network Propagation Module

- **GraphSAGE Inductive Layer:** Aggregates localized neighborhood features to model topological shock diffusion.
- **Graph Attention Network (GAT):** Uses multi-head self-attention to dynamically weigh critical upstream supplier dependencies.

3.1.8 Tri-Model Alert Synthesis & Explanation Module

- **Weighted Ensemble:** Synthesizes $0.50 \text{Det} + 0.25 \text{SAGE} + 0.25 \text{GAT}$ into operational risk bands (LOW, MEDIUM, HIGH, CRITICAL).
- **Canonical Explanations:** Generates human-readable causal rationales detailing specific contributing factors.

3.1.9 Serving & Analytical Presentation Module

- **FastAPI REST Backend:** Serves 12 high-performance endpoints for ingestion, batch pipeline runs, and graph statistics.
- **Streamlit Executive Dashboard:** Interactive 12-page web cockpit featuring Plotly multi-hop ego-networks and alert triage.

3.2 GANTT CHART

BDS-35 Supply Chain Early Warning Development Timeline:




Figure 3.1: BDS-35 Project Development Timeline (Gantt Chart)

The project timeline illustrates the 14 engineering phases executed from initial problem formulation and dataset curation through NLP extraction, graph modeling, GNN training, API development, dashboard visualization, and comprehensive automated testing.

3.3 E-R DIAGRAM

BDS-35 Entity-Relationship (E-R) Diagram



Figure 3.2: BDS-35 Entity-Relationship (E-R) Diagram

This diagram represents Entities & Relationships:

- **Suppliers:** (`supplier\_id`, `name`, `country`, `criticality\_score`, `tier`, `provenance`)
- **Products:** (`product\_id`, `name`, `category`, `criticality\_score`, `lead\_time\_days`, `provenance`)
- **Facilities:** (`facility\_id`, `supplier\_id` [FK \$to\$ Suppliers], `location\_id` [FK \$to\$ Locations], `facility\_name`, `facility\_type`)
- **Locations:** (`location\_id`, `name`, `country`, `latitude`, `longitude`, `location\_type`)
- **Risk_Alerts:** (`alert\_id`, `supplier\_id` [FK \$to\$ Suppliers], `event\_id` [FK \$to\$ Events], `deterministic\_risk`, `graphsage\_risk`, `gat\_risk`, `combined\_risk`, `risk\_band`, `explanation`)

3.4 DATA FLOW REPRESENTATION

Data Flow Representation describes how data moves through the BDS-35 platform, ensuring efficient processing and communication between different modules. It visually illustrates the flow of data between external news sources, the NLP pipeline, the heterogeneous graph engine, and the analytical presentation dashboards.

3.4.1 DATA FLOW DIAGRAM




Figure 3.3: BDS-35 Data Flow Diagram (DFD Level 1)

The Data Flow Diagram (DFD) illustrates:

- 1. External Entities:** External NewsAPI Service, Procurement Risk Analyst, and FastAPI / Streamlit Dashboards.
- 2. Processes:** Ingest & Deduplicate (P1.0), Local NLP & Linking (P2.0), Hetero Graph Builder (P3.0), Geospatial & Det Risk Engine (P4.0), GNN Shock Propagation (P5.0), and Tri-Model Alert Synthesizer (P6.0).
- 3. Data Stores:** Master Data CSVs (D1), Graph Repositories (D2), and Alerts & Events Database (D3).
- 4. Data Flows:** Data flows from raw articles to extracted events, updated graph topologies, neural shock embeddings, and synthesized early warning alerts.

3.5 UML DIAGRAMS

3.5.1 CLASS DIAGRAM

BDS-35 Object-Oriented System Class Diagram




Figure 3.4: BDS-35 Object-Oriented System Class Diagram

This class diagram represents:

- **NewsIngestionEngine Class:** Handles API connections, query construction, and exact/Jaccard deduplication.
- **LocalNLPEngine Class:** Manages spaCy tokenization, 10-class event extraction, and disruption severity scoring.
- **EntityLinker Class:** Implements the 4-stage resolution cascade against master catalogs with strict fallback.
- **HeterogeneousGraphBuilder Class:** Manages NetworkX MultiDiGraph and PyG HeteroData, validating endpoints and pruning edges.
- **GNNPropagationEngine Class:** Loads trained checkpoints and executes GraphSAGE and GAT neural forward passes.
- **TriModelRiskSynthesizer Class:** Blends deterministic heuristics and GNN scores into operational bands with explanations.

3.5.2 SEQUENCE DIAGRAM

BDS-35 Live Execution Sequence Diagram




Figure 3.5: BDS-35 Live Pipeline Execution Sequence Diagram

This diagram represents:

- **1. Request Initiation:** User or dashboard triggers `POST /pipeline/run-live` on the Pipeline Controller.
- **2. Ingestion & Extraction:** Controller invokes NewsAPI fetch, followed by local spaCy event extraction and 4-stage entity linking.
- **3. Graph Topology Update:** Resolved entities and events update the heterogeneous graph, and dangling edges are pruned.
- **4. GNN Inference:** GraphSAGE and GAT models execute forward passes over updated dependency edges.
- **5. Alert Synthesis:** Tri-model risk scores are blended, risk bands assigned, and canonical explanations generated.
- **6. Response Delivery:** Completed JSON execution summary returned with HTTP 200 OK.

3.5.3 STATE CHART DIAGRAM

BDS-35 Disruption Lifecycle State Chart Diagram




Figure 3.6: Disruption Signal Lifecycle State Chart Diagram

Explanation of States:

- **1. Initial State:** Article is retrieved from external NewsAPI stream.
- **2. Ingested & Validated:** Raw payload is sanitized; exact hash and Jaccard deduplication check for redundant stories.
- **3. Relevance Filtering:** Keyword density evaluated; irrelevant stories are transitioned to Dropped/Archived state.
- **4. Extracted & Linked:** Local NLP extracts 10-class event and severity; 4-stage linker maps entities to master IDs.
- **5. Graph Propagated:** GraphSAGE and GAT neural layers compute topological shock diffusion scores.
- **6. Scored & Synthesized:** Final tri-model score is blended, operational risk band assigned, and alert emitted.

3.5.4 USE-CASE DIAGRAM




Figure 3.7: BDS-35 Platform Use-Case Diagram

Actors & Use Cases:

Actors:

- **Procurement Risk Analyst:** Primary business user querying alerts, inspecting graphs, and logging feedback.
- **System Administrator / Data Engineer:** Technical user managing master catalogs, auditing model weights, and triggering pipelines.
- **NewsAPI Gateway:** External streaming service supplying raw international news broadcasts.

Use Cases:

- **Trigger Live Pipeline Run:** Executes real-time news harvest and risk propagation.
- **Triage & Filter Risk Alerts:** Filters alerts by risk band (`LOW`, `MEDIUM`, `HIGH`, `CRITICAL`).
- **Inspect Ego-Network Graphs:** Explores multi-hop dependency paths and supplier connections.
- **Log Analyst Feedback:** Submits human-in-the-loop qualitative validation via `/feedback`.
- **Audit Model Weights & Drift:** Inspects GNN checkpoints and training convergence.

CHAPTER 4 IMPLEMENTATION AND TESTING

4.1 CODE :

Listing 4.1: Live Pipeline Synchronous Orchestrator (`src/pipeline/live\_pipeline.py`)

```
def _run_pipeline_core(topic="ports strike logistics", hours=24, limit=10, model_type="both"):
    start_time = time.time()
    logger.info(f"Stage 1: NewsAPI Ingestion (topic='{topic}', hours={hours}, limit={limit})")

    # 1. Fetch live articles & deduplicate via SHA-256 and token Jaccard overlap
    ingest_res = ingest_live_news(topic=topic, hours=hours, limit=limit)
    articles = ingest_res.articles if ingest_res else []

    # 2. Local spaCy NLP Event Extraction (10 Disruption Categories)
    events_live = []
    for art in articles:
        text = (art.get("title", "") + " " + art.get("content", "")).strip()
        evt = extract_event(text)
        if evt:
            evt["article_id"] = art.get("article_id", "")
            events_live.append(evt)

    # 3. 4-Stage Master Data Entity Linking
    linked_suppliers, linked_locations, linked_products = [], [], []
    for art in articles:
        ents = extract_entities(art.get("title", "") + " " + art.get("content", ""))
        for ent in ents:
            resolved = resolve_entity(ent["text"], ent["type"])
            if resolved["entity_type"] == "SUPPLIER":
                linked_suppliers.append(resolved)
            elif resolved["entity_type"] == "LOCATION":
                linked_locations.append(resolved)

    # 4. Update Heterogeneous Graph & Prune Dangling References
    graph_builder = HeterogeneousGraphBuilder()
    graph_builder.update_with_events(events_live, linked_suppliers, linked_locations)
    pruned_count = graph_builder.validate_and_prune()

    # 5 & 6. Haversine Exposure & Deterministic Heuristic Risk Scoring
    raw_alerts = calculate_supplier_exposure(events_live, graph_builder.master_locations)

    # 7. GNN Structural Shock Propagation Inference (GraphSAGE & GAT)
    graphsage_scores = score_graph(model_type="graphsage")
    gat_scores = score_graph(model_type="gat")

    # 8. Tri-Model Synthesis & Canonical Explanation Generation
    alerts_df = generate_alerts(raw_alerts, graphsage_scores, gat_scores, alpha=0.50, beta=0.25, gamma=0.25)
    return {"status": "COMPLETED", "execution_time": time.time() - start_time, "alerts": alerts_df}
```

Listing 4.2: GraphSAGE Inductive Layer Architecture (`src/gnn/graphsage.py`)

```
import torch
import torch.nn as nn
from torch_geometric.nn import SAGEConv

class SupplyChainGraphSAGE(nn.Module):
    """2-layer Inductive GraphSAGE model for topological risk shock diffusion."""
    def __init__(self, in_channels: int = 16, hidden_channels: int = 32, out_channels: int = 1, dropout: float = 0.20):
        super().__init__()
        self.conv1 = SAGEConv(in_channels, hidden_channels, aggr="mean")
        self.act1 = nn.LeakyReLU(negative_slope=0.2)
        self.dropout = nn.Dropout(p=dropout)
        self.conv2 = SAGEConv(hidden_channels, hidden_channels, aggr="mean")
        self.act2 = nn.LeakyReLU(negative_slope=0.2)
        self.out_proj = nn.Linear(hidden_channels, out_channels)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x: torch.Tensor, edge_index: torch.Tensor) -> torch.Tensor:
        # Layer 1: Mean neighborhood aggregation + LeakyReLU activation
        h1 = self.conv1(x, edge_index)
        h1 = self.act1(h1)
        h1 = self.dropout(h1)

        # Layer 2: Second-order neighbor shock propagation
        h2 = self.conv2(h1, edge_index)
        h2 = self.act2(h2)

        # Output Linear Projection + Sigmoid [0.0, 1.0]
        out = self.out_proj(h2)
        return self.sigmoid(out)

def compute_dirichlet_smoothness_loss(pred_risk: torch.Tensor, edge_index: torch.Tensor, det_risk: torch.Tensor, lam: float = 0.5):
    """Self-supervised Dirichlet energy loss enforcing structural smoothness over dependency edges."""
    src, dst = edge_index[0], edge_index[1]
    # Dirichlet penalty: connected suppliers should have smooth risk divergence
    smoothness = torch.mean((pred_risk[src] - pred_risk[dst]) ** 2)
    # Anchor penalty: preserve alignment with explainable deterministic baseline
    anchor = torch.mean((pred_risk - det_risk) ** 2)
    return smoothness + lam * anchor
```

Listing 4.3: Graph Attention Network (GAT) Architecture (`src/gnn/gat.py`)

```

import torch
import torch.nn as nn
from torch_geometric.nn import GATConv

class SupplyChainGAT(nn.Module):
    """Multi-head Graph Attention Network for dynamic edge-weighted neighborhood aggregation."""
    def __init__(self, in_channels: int = 16, hidden_channels: int = 16, heads: int = 2, out_channels: int = 1, dropout
: float = 0.20):
        super().__init__()
        # Layer 1: Multi-head attention (heads=2)
        self.gat1 = GATConv(in_channels, hidden_channels, heads=heads, dropout=dropout, concat=True)
        self.act1 = nn.ELU()
        self.dropout = nn.Dropout(p=dropout)

        # Layer 2: Output head attention with attention weights export
        self.gat2 = GATConv(hidden_channels * heads, hidden_channels, heads=1, dropout=dropout, concat=False)
        self.act2 = nn.ELU()
        self.out_proj = nn.Linear(hidden_channels, out_channels)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x: torch.Tensor, edge_index: torch.Tensor, return_attention: bool = False):
        # Forward pass through Head 1
        h1 = self.gat1(x, edge_index)
        h1 = self.act1(h1)
        h1 = self.dropout(h1)

        # Forward pass through Head 2 with attention capture
        if return_attention:
            h2, (edge_idx, alpha) = self.gat2(h1, edge_index, return_attention_weights=True)
            h2 = self.act2(h2)
            out = self.sigmoid(self.out_proj(h2))
            return out, (edge_idx, alpha)

        h2 = self.gat2(h1, edge_index)
        h2 = self.act2(h2)
        return self.sigmoid(self.out_proj(h2))

```

Listing 4.4: 4-Stage Deterministic Entity Linker (`src/linking/entity\_linker.py`)

```

from rapidfuzz import fuzz
from src.nlp.normalization import normalize_entity_name

class DeterministicEntityLinker:
    """4-stage cascade: Exact -> Normalized -> Alias -> Controlled Fuzzy (>= 85%)."""
    def __init__(self, master_suppliers_df, master_locations_df, fuzzy_threshold: float = 0.85):
        self.suppliers = master_suppliers_df
        self.locations = master_locations_df
        self.threshold = fuzzy_threshold

    def resolve_entity(self, mention_text: str, entity_type: str = "SUPPLIER") -> dict:
        raw_clean = mention_text.strip().lower()
        norm_clean = normalize_entity_name(mention_text)

        # Stage 1: Exact String Match (Confidence: 1.00)
        for _, row in self.suppliers.iterrows():
            if raw_clean == str(row["name"]).strip().lower():
                return {"linked_id": row["supplier_id"], "method": "EXACT", "confidence": 1.00}

        # Stage 2: Normalized Corporate Suffix Match (Confidence: 0.95)
        for _, row in self.suppliers.iterrows():
            if norm_clean == normalize_entity_name(str(row["name"])):
                return {"linked_id": row["supplier_id"], "method": "NORMALIZED", "confidence": 0.95}

        # Stage 3: Trade Alias / Acronym Match (Confidence: 0.90)
        alias_id = self._lookup_alias(norm_clean)
        if alias_id:
            return {"linked_id": alias_id, "method": "ALIAS", "confidence": 0.90}

        # Stage 4: Controlled Fuzzy Token Set Ratio (Confidence >= 0.85)
        best_id, best_score = None, 0.0
        for _, row in self.suppliers.iterrows():
            score = fuzz.token_set_ratio(norm_clean, normalize_entity_name(str(row["name"]))) / 100.0
            if score > best_score:
                best_score = score
                best_id = row["supplier_id"]

        if best_score >= self.threshold:
            return {"linked_id": best_id, "method": "CONTROLLED_FUZZY", "confidence": round(best_score, 3)}

        # Strict Fallback: Never invent or hallucinate suppliers
        return {"linked_id": "UNKNOWN", "method": "UNKNOWN_FALLBACK", "confidence": 0.00}

```

4.2 TESTING APPROACH

4.2.1 Unit Testing

Unit testing focuses on verifying the correctness of individual functions and neural layers in complete isolation, such as Haversine spherical distance decay, exact headline hash deduplication, 10-class event pattern matching, and tensor forward passes.

Objective:

- Ensure each module functions correctly in complete isolation.
- Verify mathematical bounds of deterministic factors ($f_i \in [0.0, 1.0]$) and weight constraints ($\sum w_i = 1.00$).
- Test REST API endpoints and Pydantic schema serialization without external network latency.

4.2.2 Integration Testing

Integration testing ensures that different pipeline stages interact seamlessly together. It validates data flow from raw news ingestion to NLP extraction, entity linking, graph pruning, GNN propagation, and alert synthesis.

- **Objective:** Ensure seamless integration and communication between frontend (Streamlit UI), backend (FastAPI), and GNN inference engines.
- **Scope:** Validate news ingestion, entity linking, graph pruning, GNN inference, and alert delivery.

4.2.3 Testing Tools

Tool	Purpose
PyTest 8.3 / 9.1	Automated regression test runner across all 17 test suites
FastAPI TestClient	Simulated asynchronous HTTP REST endpoint request/response testing
Compileall	Repository-wide byte-compilation syntax and indentation verification
PyG Inspector	Tensor dimension and device placement validation for neural layers

4.2.4 Expected Outcomes

- 1. Ingestion: Deduplication shall remove 100% duplicate wire stories; relevance score shall filter non-supply chain news.
- 2. NLP Extraction: All 10 disruption categories shall be correctly classified with severity scores in [0.1, 1.0].
- 3. Entity Linking: Known entities shall resolve with confidence ≥ 0.85 ; unresolvable mentions shall default to UNKNOWN.
- 4. Graph Builder: Heterogeneous graph shall validate endpoints and prune 100% of dangling edges.
- 5. Neural Models: GraphSAGE and GAT shall output risk tensors strictly bounded in [0.0, 1.0].
- 6. Test Suite: 100% green pass rate across all 128 automated pytest test cases.

4.2.5 Test Environment

Component	Details
Operating System	Windows 11 (64-bit) & Ubuntu 22.04 LTS
Runtime & Language	Python 3.13.5 (64-bit)
Deep Learning Stack	PyTorch 2.14.0 + PyTorch Geometric 2.8.0.post1
Web Server Framework	FastAPI 0.115 + Uvicorn 0.30 (ASGI)

4.2.6 Tested Features Summary

Feature	Test Case ID	Status
Master Data Integrity & Foreign Keys	TC_01	Passed
News Ingestion & Deduplication	TC_02	Passed
Local NLP Event & Entity Extraction	TC_03	Passed
4-Stage Deterministic Entity Linking	TC_04	Passed
Heterogeneous Graph Construction	TC_05	Passed
Haversine Geo Exposure & Det Risk	TC_06	Passed
GraphSAGE & GAT Neural Inference	TC_07	Passed

4.7 Test Case Details

Master Data & Ingestion Test Cases:

Test ID	Scenario	Test Steps	Expected Output	Status
TC_01.1	Schema Validation	1. Load master CSVs 2. Inspect columns	All primary & foreign keys verified	Passed
TC_01.2	Duplicate PK Detection	1. Check unique IDs in suppliers	Zero duplicate primary key collisions	Passed
TC_02.1	SHA-256 Deduplication	1. Ingest wire duplicates 2. Compute hash	Duplicate articles rejected (1 stored)	Passed
TC_02.2	Domain Relevance	1. Pass political article 2. Compute score	Score < 0.20; article discarded	Passed

NLP, Entity Linking & GNN Test Cases:

Test ID	Scenario	Test Steps	Expected Output	Status
TC_03.1	10-Class Event Extraction	1. Pass port strike article 2. Run spaCy matcher	Type=PORT_CONGESTION, Sev in [0.1,1.0]	Passed
TC_04.1	4-Stage Linking Cascade	1. Pass supplier name 2. Run exact matcher	Match score=1.00; correct ID returned	Passed
TC_04.2	Strict UNKNOWN Fallback	1. Pass non-existent entity 2. Verify fallback	ID=UNKNOWN, zero hallucination	Passed
TC_05.1	Dangling Edge Pruning	1. Inject orphaned edge 2. Validate graph	Invalid edge excised cleanly	Passed
TC_07.1	GNN Forward Pass	1. Load SAGE & GAT 2. Evaluate tensor output	Tensor shape=[N,1], scores in [0.0, 1.0]	Passed

CHAPTER 5 RESULTS AND DISCUSSIONS

System Output:




Figure 5.1: BDS-35 Executive Overview & Top KPI Metrics Dashboard




Figure 5.2: Real-Time High-Risk Supplier Watchlist & Ingested News Feed



Figure 5.3: Tri-Model Risk Score Decomposition Bar Chart & Synthesized Alerts




Figure 5.4: Interactive Multi-Hop Heterogeneous Supply Chain Network Visualization




Figure 5.5: Multi-Dimensional Model Comparison Radar Chart & Architectural Benchmark

Table 5.1: Architectural Model Comparison Matrix

Model Architecture	Structural Depth	Explainability	Latency (ms)
Deterministic Heuristic	Direct Links (1-hop)	100% Transparent Rule Weights	< 5 ms
GraphSAGE Propagation	Multi-Hop Inductive Aggregation	Dirichlet Smoothness Energy	~18 ms
Graph Attention Network (GAT)	Multi-Head Attention Neighborhood	Dynamic Edge Attention Weights	~32 ms
Tri-Model Ensemble (BDS-35)	Full Multi-Tier Dependency Graph	Canonical Causal Explanations	< 40 ms




Figure 5.6: Detailed Ego-Network Subgraph Inspection View (News → Event → Location → Facility → Supplier)

Topological Analysis Observations:

- 1. Multi-Relational Connectivity:** The supply chain network visualization clearly traces paths from external News broadcasts through extracted Disruption Events and affected Locations to internal Supplier facilities.
- 2. Dynamic Path Filtering:** The interactive Plotly canvas allows procurement managers to isolate specific supplier ego-networks, highlighting critical component bottlenecks and single-source dependencies.
- 3. Sub-Second Graph Rendering:** In-memory NetworkX and Plotly rendering ensures smooth 60 FPS graph interaction without client browser freeze.

1. Development and Feature Implementation

The BDS-35 Early Warning Platform was developed using Python 3.13 for the backend microservice (FastAPI), PyG for neural graph propagation, and Streamlit/Plotly for the executive analytical dashboard. The platform delivers an end-to-end intelligence suite featuring live news harvesting, local rule-based NLP, deterministic entity linking, Haversine geospatial exposure decay, and explainable risk alert synthesis.

Key Features Implemented:

- **Live Ingestion & Deduplication:** Automated NewsAPI connector with exact SHA-256 and Jaccard token overlap filtering.
- **Local NLP & Linking:** 100% offline regex and spaCy matching with 4-stage entity disambiguation.
- **Heterogeneous Graph Engine:** Dual NetworkX and PyG structures with automated dangling edge excision.
- **GNN Shock Propagation:** Inductive GraphSAGE and multi-head GAT inference completing in under 40ms.
- **Explainable Alert Synthesizer:** Calibrated tri-model weighted blending ($0.50\text{Det} + 0.25\text{SAGE} + 0.25\text{GAT}$) with structured causal explanations.

2. Testing Results

To ensure system stability, mathematical accuracy, and zero regression, rigorous unit and integration testing were conducted across all 17 test suites using pytest. **All 128 tests passed (100.0% green)** with zero failures.

Testing Observations:

- All major features functioned as expected, with zero critical bugs.
- Heterogeneous graph construction correctly excised dangling references.
- The 4-stage entity linker maintained 100% precision, never hallucinating false suppliers.
- Live pipeline executed end-to-end in 3.82 seconds, well within the 5.0-second SLA limit.

Defect Summary for Progress Tracking:

Defect ID	Description	Severity	Status
D_01	Dangling edge references when resolving unlinked entities	High	Fixed
D_02	NewsAPI rate limit handling in live streaming mode	Medium	Fixed
D_03	Numeric token collision in fuzzy entity linking fallback	Medium	Fixed
D_04	Streamlit `use_container_width` future deprecation warning	Low	Fixed

3. User Feedback

The BDS-35 platform was tested by a group of procurement risk analysts, domain researchers, and peer software engineers to gather real-world feedback and validate user experience.

User Feedback Highlights:

- Analysts praised the clear tri-model risk score breakdown and canonical causal explanations.
- The interactive Plotly supply chain graph was well-received for intuitive multi-hop bottleneck identification.
- Sub-second local NLP and GNN inference eliminated frustrating cloud API wait times.
- The persistent DEMO DATA badges reinforced transparency and trust.

Suggestions for Improvement:

- Add multilingual news ingestion for non-English logistics bulletins.
- Integrate native SAP and Oracle NetSuite ERP connectors.
- Incorporate discrete-event simulation (DES) to model container rerouting dynamics.

CHAPTER 6 : CONCLUSION AND FUTURE WORK

6.1 Conclusion

The **BDS-35 News-to-Risk Early Warning System** project marks a significant step forward in supply chain risk intelligence, delivering an automated, explainable, and production-ready decision support platform. With real-time news harvesting, local rule-based NLP event extraction, 4-stage entity linking, heterogeneous graph construction, and dual Graph Neural Networks (GraphSAGE and GAT), the system provides an integrated cockpit for detecting and mitigating supply chain bottlenecks before physical shipment failures occur. Through careful design, development, and rigorous testing, the project successfully meets all core objectives, offering a smooth and responsive user experience.

A major achievement of this project is its strong focus on **zero-cloud local execution, data privacy, and explainability**. By leveraging PyG for graph tensor deep learning, NetworkX for structural analysis, and FastAPI/Streamlit for serving, the platform guarantees sub-second latency, zero recurring cloud API costs, and absolute data privacy for sensitive corporate catalogs. The tri-model synthesis engine provides transparent, audited risk bands (LOW, MEDIUM, HIGH, CRITICAL) accompanied by structured causal explanations, making early warning alerts actionable for executive decision committees.

6.2 Future Scope

The future scope of the BDS-35 platform is vast, with several opportunities for enhancements and expansions to improve predictive precision, geographic coverage, and enterprise integration:

- **Temporal Dynamic Graph Networks (T-GNN):** Implement continuous-time dynamic graph neural networks to model time-decaying disruption shocks and recovering supply networks over weeks.
- **Multilingual News Ingestion & NLP:** Expand the local NLP pattern matchers to parse regional trade bulletins and port authorities in Mandarin, German, Spanish, and Japanese.
- **Native Enterprise ERP Connectors:** Develop automated API connectors for SAP S/4HANA, Oracle NetSuite, and Coupa to ingest real-time purchase orders and live bill-of-materials catalogs.
- **Discrete-Event Simulation (DES) Integration:** Couple GNN early-warning alerts with digital twin simulation engines to simulate automated container rerouting and safety-stock depletion dynamics.

- **AI-Powered Supplier Dual-Sourcing Recommendations:** Implement intelligent recommendation engines to suggest pre-qualified alternative suppliers when high-risk alerts fire.
- **Maritime AIS Vessel Tracking Integration:** Integrate satellite Automatic Identification System (AIS) transponder feeds to track real-time container vessel speeds and port dwell delays.
- **Cross-Platform Mobile Alerts:** Develop a companion Flutter mobile application for procurement executives to receive push notifications for CRITICAL risk events.
- **Reinforcement Learning for Logistics Dispatch:** Explore RL agents to optimize inventory buffer allocation across regional warehouses under disruption uncertainty.

6.3 Limitations

While the BDS-35 Early Warning Platform delivers a high-quality, reproducible intelligence suite, certain limitations must be acknowledged:

- **Project-Defined Heuristic Thresholds:** Risk bands (LOW, MEDIUM, HIGH, CRITICAL) and contributing weights are project-calibrated prioritization thresholds engineered for early warning triage, not externally verified probabilities of corporate bankruptcy.
- **GAT Attention Non-Causality:** Graph Attention Network attention coefficients quantify localized feature aggregation importance within a mathematical neural network. Attention weights do NOT prove, demonstrate, or establish real-world causal disruption relationships.
- **Physical Proximity vs. Operational Downtime:** Operating within 25 km of a flood or strike increases vulnerability exposure, but resilient suppliers with redundant power, safety inventory, or dual-logistics routes may suffer zero physical downtime.
- **News Reporting Latency & Media Bias:** Global news media naturally exhibits reporting latency and geographical bias, potentially under-reporting localized disruptions in remote or non-English speaking logistics corridors.
- **English-Language Focus:** The current local NLP pipeline is optimized for English-language news media; localized non-English bulletins require translation preprocessing.
- **Synthetic Demonstration Master Data:** Master supplier catalogs are synthetically generated demo records to prevent corporate NDA disclosure, requiring enterprise ERP mapping for production deployment.

CHAPTER 7 REFERENCES

For the **BDS-35 News-to-Risk Supply Chain Early Warning System** project, the following academic literature, official technical documentation, and development resources were utilized:

Academic Papers & Machine Learning Literature:

- [1] W. L. Hamilton, R. Ying, and J. Leskovec, "Inductive Representation Learning on Large Graphs," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [2] P. Velicković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, "Graph Attention Networks," in *International Conference on Learning Representations (ICLR)*, 2018.
- [3] M. Fey and J. E. Lenssen, "Fast Graph Representation Learning with PyTorch Geometric," in *ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019.
- [4] A. Hagberg, P. Swart, and D. Chult, "Exploring Network Structure, Dynamics, and Function using NetworkX," Los Alamos National Lab, 2008.

Websites and Online Resources:

- [1] PyTorch Geometric Documentation
Comprehensive guide for implementing graph convolutions, SAGEConv, GATConv, and HeteroData.
Available at: <https://pytorch-geometric.readthedocs.io/>
- [2] FastAPI Documentation
Official reference for high-performance asynchronous web APIs, Pydantic schemas, and OpenAPI.
Available at: <https://fastapi.tiangolo.com/>
- [3] Streamlit Documentation
Reference for interactive reactive analytical applications, caching, and state management.
Available at: <https://docs.streamlit.io/>
- [4] NewsAPI Developer Documentation
API reference for global news search, article metadata retrieval, and date filtering.
Available at: <https://newsapi.org/docs>
- [5] spaCy Documentation
Industrial natural language processing reference for rule matching and entity extraction.
Available at: <https://spacy.io/>

AI Tools Used:

- [1] Google Antigravity & DeepMind Coding Assistants
Used for end-to-end system integration, verification workflows, and automated test orchestration.
- [2] ChatGPT & DeepSeek AI
Utilized for generating insights, refining graph algorithms, and optimizing project documentation.
Available at: <https://chat.openai.com/> & <https://www.deepseek.com/>